Benchmarking Faithfulness of Mathematical Chain-of-Thought Autoformalisation

Anonymous Authors[1]

Abstract

Autoformalisation (AF) of chain-of-thought (CoT) is increasingly used to verify mathematical reasoning by translating natural-language steps into a formal language such as Lean. Yet, current evaluations largely reward *validity preservation*—formalising true statements into provable goals—while overlooking the complementary requirement of *invalidity preservation*: faithfully formalising false statements into goals that remain unprovable (e.g., avoiding "sycophantic" auto-correction). We propose a bifaceted notion of AF faithfulness with the above positive and negative aspects, and introduce a benchmarked metric that balances both. Our method automatically perturbs correct CoT statements to inject controlled errors, then measures (i) proof success on unperturbed steps and (ii) proof failure on perturbed steps, enabling scalable application of our method to diverse CoT datasets. Across state-of-the-art AFs and multiple mathematical CoT datasets, we observe pervasive sycophancy: many AFs silently repair false inputs into provable statements. The most validity-preserving fine-tuned AFs are also the most sycophantic. Our empirical results suggest a tension between validity and invalidity preservation, motivating holistic benchmarks (generated by our method) for faithfully verifying mathematical CoTs.

1. Introduction

Neural Theorem Proving (NTP) applies Large Language Models (LLMs) and other Machine Learning techniques to construct mathematical proofs that are formalised in and verified by proof assistants like Lean (de Moura & Ullrich, 2021), Rocq (Coquand & Huet, 1988), and Isabelle (Paulson, 1990). Numerous recent efforts (Xin et al., 2025; Ren et al., 2025; Lin et al., 2025; Chen et al., 2025; Xu et al., 2025; Wang et al., 2025a) have yielded promising results, achieving gold-medal-level performance at International Mathematical Olympiad 2025 (Achim et al., 2025; Luong & Lockhart) and near-saturation results on miniF2F (Zheng et al., 2022) and PutnamBench (Tsoukalas et al., 2024). Beyond benchmarking competitions, NTP is making its mark in practical applications. Recent works (Quan et al., 2025; Zhou & Zhang, 2025; Liu et al., 2025a) have attempted to use NTP to automatically verify the correctness of LLM-generated Chain-of-Thought (CoT) in mathematical reasoning. This is the setting of our paper.

Limitations of Prior Work Despite the progress, a fundamental difficulty persists in using NTP to verify mathematical CoT: proof assistants verify statements in formal languages, whereas LLMs reason in natural languages. This gap necessitates so-called *autoformalisation* (AF) to translate natural-language CoTs into formal statements. However, dominant AF solutions still rely on LLMs, leaving the *faithfulness* of this translation unguaranteed: it is difficult to determine whether the semantics of the original natural language statement is preserved in the translated formal statement.

Various methods have been proposed to improve AF faithfulness, such as using training datasets constructed with extensive human-expert supervision (Gao et al., 2024b; Wang et al., 2025b; Zheng et al., 2025), and more rigorous evaluation by comparing AF model outputs against human-labeled ground truth using various metrics (Liu et al., 2025b;c). However, these methods typically assume that the given natural language statements are correct, and require AF to produce semantics-preserving formalisations of these correct statements. Alas, this assumption does not hold in the CoT verification setting – the very reason we verify CoT is that it may be incorrect! Indeed, we find that no prior work has systematically examined whether existing AF models can preserve errors in incorrect natural language statements and faithfully represent these errors in the resulting formalisations.

Our Approach To investigate this, we deliberately perturb correct CoTs (selected from public datasets) by intro-

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

ducing errors and examining whether AF models can faithfully preserve these errors. Our experimental results are alarming. Many state-of-the-art AF models—specifically, LLMs fine-tuned to translate natural language CoTs into a formal language—exhibit high levels of sycophancy: they silently correct the deliberately introduced errors, producing corrected statements that should not be corrected. This can lead to a serious problem: when a CoT contains errors due to LLM hallucination, the AF model in a verification pipeline may silently correct the hallucinated error, passing the corrected formal statement to the NTP backend, causing the flawed CoT to be incorrectly verified, thereby invalidating the whole verification.

Example 1.1 (*Ex Falso Quodlibet* via Impossible Constraint). **Problem ID:** `omnimath-629`. The original problem assumes that x is a real number. The perturbed content asserts that $3^{(x+2)} = 51$. The AF Kimina silently changes the type of x to $\mathbb{N}$, and hence, no value exists that satisfies the equation (which is a premise). Consequently, the conclusion follows from a contradiction.

```
theorem my_favorite_theorem
  (x : \mathbb{N})
  (h : 3^x = 5) :
  3^(x + 2) = 51 := by
    have h : False := by...
```

See Appendix F.1 for more examples of such AF sycophantic behaviour.

Contributions Motivated by this potential critical issue, we make the following contributions.

1. We propose a definition of AF faithfulness that emphasises proper semantic preservation, even when the input contains errors, and an accompanying faithfulness metric called *faithfulness index*.

2. By creating our own ground truth of "wrongness", we introduce a scalable, model-agnostic, perturbation-based methodology to automatically generate faithfulness benchmarks.

3. We present *FaithformBench*, a benchmark derived from applying our method.

On this benchmark, we evaluate 4 models fine-tuned specifically for the AF task: Goedel (Lin et al., 2025), Herald (Gao et al., 2024b), Kimina (Wang et al., 2025a), and Stepfun-Formaliser (Wu et al., 2025); and 4 general-purpose foundation models (DeepSeek-V3.2, GPT-5.2, Gemini-Pro, and Qwen-3-Max), arriving at the following findings:

1. All the fine-tuned models exhibit high levels of sycophancy, correcting errors when they should not, rather than faithfully representing them.

2. The more capable a fine-tuned model is in formalising correct CoTs, the more sycophantic it is (i.e., the more likely it is to silently correct errors).

3. Compared to the fine-tuned models, the general-purpose models exhibit a significantly lower sycophantic behaviour in end-to-end formalisation of CoT into Lean proofs.

To the best of our knowledge, this is the first benchmark designed to assess the *faithfulness* of autoformalisation in the context of verifying mathematical CoT. Moreover, unlike prior approaches such as BEq (Liu et al., 2025b) and GTED (Liu et al., 2025c), our method enables automatic dataset construction without requiring human-expert-labeled ground truth. Our contributions provide a diagnostic tool for the community, highlighting that before AF can be trusted to guide CoT reasoning, we must solve the problems of formalising real-world natural language statements and autoformalisation sycophancy.

2. Related Work

Neural Theorem Proving and Formalisation Neural theorem proving is a rapidly evolving field. LeanDojo (Yang et al., 2023) introduces a retrieval-augmented generation (RAG) framework. LEGO-Prover (Wang et al., 2024a) emphasizes the automated synthesis and reuse of lemmas as LLM skills. Draft, Sketch, and Prove (Jiang et al., 2023) proposes that NTP models can first generate natural language proof drafts, then formalize them into formal proofs. DeepSeek-Prover-V1.5 (Xin et al., 2025) applies Reinforcement Learning and Monte-Carlo Tree Search. DeepSeek-Prover-V2 (Ren et al., 2025), Seed Prover (Chen et al., 2025), and Kimina (Wang et al., 2025b) further incorporate CoT in their reasoning.

Autoformalisation and Formal Verification in the Wild Wu et al. (2022) set the stage by bridging natural language and machine-verifiable code, while Zhou et al. (2024) sharpened this mandate by grounding quantitative reasoning through autoformalised checks. Recent efforts focus on making verification more granular: Herald (Gao et al., 2024b) was an early work that released its dataset and formaliser. This was followed by Zhou & Zhang (2025) and Liu et al. (2025a), which utilise step-by-step formal feedback to catch reasoning errors mid-solution, effectively acting as a running companion rather than a post-hoc audit. ProofBridge (Jana et al., 2025) aligns natural language and formal proofs via shared semantic embeddings,

while StepFun-Formalizer (Wu et al., 2025) fuses knowledge with reasoning.

Sycophancy in LLM Sycophancy (Perez et al., 2023; Sharma et al., 2024) in general refers to an LLM's tendency to agree with the user's opinions or preferences, even if they are false or unethical. In the context of mathematics, a strand of work (Liu et al., 2026; Kirichenko et al., 2025; Xue et al., 2025) has been devoted to studying how an LLM will answer an ill-formed (underspecified or unreasonable) mathematical question. BrokenMath (Petrov et al., 2025) is closest to our work: they study sycophancy in (natural language) theorem proving, where an LLM might come up with a hallucinated proof in natural language for an incorrect statement. In formalised theorem proving, that problem is not too harmful because the proof will simply be rejected by the proof checker. In this paper, we identify and study a novel type of sycophancy that arises in autoformalisers, specifically, LLMs fine-tuned to translate natural language statements into Lean representations such as Goedel, StepFun, Kimina and Herald.

Data, Benchmarks, and Evaluation Metrics The rapid scaling of autoformalisation models has necessitated more rigorous benchmarks and nuanced evaluation metrics. To this end, datasets such as Herald (Gao et al., 2024b) provide essential natural-language-annotated Lean 4 data. However, evaluating formalisation on these complex datasets goes beyond mere binary correctness; it requires assessing the structural alignment and faithfulness of the generated code. Recent metrics offer more precise, semantically-grounded measurements of this formalisation quality: Generalized Tree Edit Distance (GTED) (Liu et al., 2025c) quantifies structural similarity, while Bidirectional Extended Definitional Equivalence (BEq) (Liu et al., 2025b) leverages formal-grounded equivalence to measure accuracy. For the CoT process, ProcessBench (Zheng et al., 2025) audits steps in natural language; we audit steps in formal language, building upon this line of inquiry.

3. Problem Setup and Benchmark Definition

Modelling CoT For our purposes, a *chain-of-thought (CoT)* is simply a string x that contains a reasoning trace. We are deliberately loose about the "encoding" used for x, which can vary significantly between datasets. For example, x may consist of unstructured text describing a "raw" reasoning trace, or it may consist of a JSON object with certain reasoning "steps" already delineated (as in Zheng et al. (2025)).

Importantly, we do not assume that the input string x always corresponds to a *valid* reasoning trace. This is because LLMs (and humans) frequently make logical mistakes, and so it is very common to encounter inconsistent CoTs in real datasets.

AF output We are interested in AF systems that map an input CoT x to a sequence of formal statements as follows:

$$\Gamma_1 \vdash A_1 \qquad \Gamma_2 \vdash A_2 \qquad \cdots \qquad \Gamma_n \vdash A_n. \tag{1}$$

Each $\Gamma_i = (\Gamma_i^1, \ldots, \Gamma_i^{n_i})$ is a tuple of *premises*, and each A_i is a *conclusion*. Every Γ_i^j and A_k is encoded in some common formal system of choosing (e.g. Lean).

The idea is that each statement $\Gamma_i \vdash A_i$ asserts that A_i can be proven from Γ_i using the inference rules of the formal system in question. In Lean, $\Gamma_i \vdash A_i$ would be encoded as follows:

```
theorem _ (p_1 : Gamma_i_1)
    ... (p_ni : Gamma_i_ni) : A_i :=
        by sorry
```
(2)

Notice that, as is often the case for AF tasks, we do not require the AF system to produce a *proof* of this theorem: only a formal statement of the theorem itself. (Also notice that (2) is well-typed in Lean, despite containing `by sorry`, which skips the proof.)

In practice, many AF methods (including the examples we consider in our experiments) are designed only to produce a single output

$$\Gamma \vdash A. \tag{3}$$

However, the approach we describe here applies equally well to the general case of (1) too.

Faithfulness We now consider how we would like our AF system ideally to behave. Intuitively, we want the output (1) to correspond semantically to the same logical argument that the input x describes. When this happens, we say the system is *faithful*. This condition seems to underpin most AF methods we are aware of in practice (in some cases implicitly).

Example 3.1. Consider the following CoT:

> "It holds that $2 > 1$. We also know that $3 > 2$. Therefore $3 > 1$ also."

It is clear that the desired output is as follows:

$$\vdash 2 > 1 \qquad \vdash 3 > 2 \qquad 2 > 1, 3 > 2 \vdash 3 > 1.$$

Now consider the following modification to the CoT:

> "It holds that $2 > 1$. We also know that $3 > 2$. Therefore $0 > 1$ also."

This CoT is invalid: in particular, its final step is clearly false. The output of a faithful AF system should also capture these errors as follows:

$$\vdash 2 > 1 \qquad \vdash 3 > 2 \qquad 2 > 1, 3 > 2 \vdash 0 > 1$$

Consequences of faithfulness As Example 3.1 illustrates, a faithful AF system must satisfy the following two conditions for each input CoT x:

- *Validity preservation:* If x is valid, then every $\Gamma_i \vdash A_i$ should be true.
- *Invalidity preservation:* If x is invalid, then some $\Gamma_i \vdash A_i$ should be false.

On their own, these conditions are not *sufficient* for faithfulness. For example, the following AF system satisfies both conditions:

$$x \mapsto \begin{cases} \vdash \text{true} & \text{if } x \text{ is valid} \\ \vdash \text{false} & \text{if } x \text{ is invalid} \end{cases} \tag{4}$$

This would clearly not be faithful, since the output throws away almost all the information contained in x. More generally, if an AF method is *not* faithful, then one of the following failure modes must hold for some CoT x:

1. *Error-inducing:* x is valid, but some $\Gamma_i \vdash A_i$ is false
2. *Sycophancy:* x is invalid, but every $\Gamma_i \vdash A_i$ is true
3. *Semantic drift:* neither of the previous two cases holds, but the output still does not correspond semantically to x in some way.

The AF system (4) suffers from the third failure mode here.

3.1. Benchmarking faithfulness

We would like a benchmark that assesses the degree to which an AF method is faithful. One approach to doing so is to take a large dataset of CoTs x, pass each through the AF system, and determine how often the output corresponds semantically to x. Unfortunately, the most natural way of doing this requires some form of ground truth formalisation for each input x, which is often too expensive to obtain at scale.

As an alternative approach, we therefore propose a benchmark that assesses only validity and invalidity preservation. Equivalently, we assess the extent to which the system is error-inducing and sycophantic. We describe below a general method to doing so based on automatically *perturbing* correct CoT inputs x. Our approach is lightweight and scalable, and in particular avoids the requirement for hand-generated ground-truth formalisations.

Before proceeding, we emphasise that our benchmark is not intended to be a *complete* test for faithfulness. As previously discussed, even if an AF system preserves validity and invalidity (or equivalently is neither error-inducing nor sycophantic), it does not follow that the system is faithful. On the other hand, under mild assumptions, our benchmark provides *sound* information about the performance of the AF system. For example, if we *do* discover failure modes, then we have learned something concrete about its performance. Our experiments show that this leads to interesting insights into the behaviour of real AF systems (which in particular exhibit a large degree of sycophancy in many cases).

Dataset assumptions We assume access to a large dataset of CoTs x. The only assumption we make about each x is that it contains a valid reasoning trace. We do not assume that x is accompanied by a ground truth formalisation. In our experiments, we obtain our data from ProcessBench (Zheng et al., 2025), although other sources could be used as a drop-in replacement.

Benchmarking validity-preservation It is reasonably straightforward to assess whether an AF system is validity-preserving using the dataset we have assumed. Since we know each CoT x contains a valid reasoning trace, we can simply:

1. Formalise each x
2. Try to determine whether the result is provable.

At a high level, if a large percentage of the examples are in fact provable, then we can conclude that the system is usually validity-preserving.

Benchmarking invalidity-preservation A key novelty of our metric lies in our approach to assessing whether the system preserves invalidity instead. For this, we proceed as follows:

1. Pass each x using a *perturbation function* (see below), obtaining a new CoT x'
2. Formalise each x'
3. Try to determine whether the result is provable.

For the perturbation functions we use, our key working assumption is that x' contains an *invalid* reasoning trace (at least for most of the examples in the dataset). Accordingly, if a large number of the x' values are in fact provable, then we can conclude that the system is often sycophantic (i.e. does not preserve invalidity). Intuitively, this indicates that the AF system is *silently correcting* x' so that it becomes provable, rather than formalising its actual semantic content.

Details Let F denote the AF method. This takes a CoT x as input and tries to produce an output $F(x)$ of the form (1). We also allow for the possibility that F fails to produce an output (case AC-FAIL), or produces an output that is syntactically invalid (e.g. an output that is not valid Lean code) (case TC-FAIL). We denote this case as $F(x) = \bot$.

In addition, let P be a *prover*. In other words, P takes unproven formal statement s like (2) as input, and tries to find a corresponding proof $P(s)$ as output. We write $P(s) = \top$ when the prover succeeds, and $P(s) = \bot$ otherwise, e.g. due to failure of the Lean kernel check, or the proof search timing out. (It is convenient to let $P(\bot) = \bot$ also.)

Finally, we let $\mathsf{Perturb}$ be a *perturbation function*. This takes as input a CoT x and produces another CoT $\mathsf{Perturb}(x)$ as output by modifying x in some way (potentially stochastically). Different choices are possible here (see below).

We then proceed as follows. Assume we have N (correct) CoTs x_i in our dataset. We then let:

$$\mathrm{TPR} := \frac{1}{N}\sum_{i=1}^{N} \mathbb{I}(P(F(x_i)) = \top)$$

$$\mathrm{FPR} := \frac{1}{N}\sum_{i=1}^{N} \mathbb{I}(P(F(\mathsf{Perturb}(x_i))) = \top)$$

Here $\mathbb{I}$ denotes an indicator function, and TPR and FPR stand for "true positive rate" and "false positive rate" respectively. Both numbers here give concrete information about the performance of the model:

- If $\mathrm{TPR} \approx 1$, then we can be confident that the AF method usually preserves validity. This is because we have assumed each x_i is logically correct, and because if $P(F(x_i)) = \top$, then $F(x_i)$ is true (since we trust the prover).
- If $\mathrm{FPR} \approx 1$, we can infer that the AF method is often sycophantic. This is because we are assuming $\mathsf{Perturb}(x_i)$ is almost always *incorrect*. Accordingly, if $F(\mathsf{Perturb}(x_i))$ is provable, then the formaliser must have changed its logical meaning in some way.

It is desirable for an AF system to have TPR close to 1, while at the same time FPR close to 0. Based on this, we further propose to aggregate these quantities in a *Faithfulness Index* (FI) defined as

$$\mathrm{FI} = \frac{2\,\mathrm{TPR}\cdot(1-\mathrm{FPR})}{\mathrm{TPR}+(1-\mathrm{FPR})}$$

In other words, FI is the harmonic mean of TPR and $1-$ FPR. This quantity becomes largest for $\mathrm{TPR} \approx 1$ and $\mathrm{FPR} \approx 0$, as desired.

Additional quantities of interest Beyond TPR and FNR, other quantities of interest may be computed using similar ideas. For example, we could define the "true negative rate" as

$$\mathrm{TNR} := \frac{1}{N}\sum_{i=1}^{N} \mathbb{I}(P(\neg F(x_i)) = \top),$$

where $\neg t$ is the logical negation of a formal statement t (where also $\neg\bot = \bot$). This quantity is large if, for many perturbed CoTs, the formaliser produces an output that is provable (since its negation is provably true). Likewise, we could define the "false negative rate" as

$$\mathrm{FNR} := \frac{1}{N}\sum_{i=1}^{N} \mathbb{I}(F(x_i) = \bot \text{ or } P(\neg F(x_i)) = \top).$$

This quantity is large if for many inputs, the formaliser either 1) fails to produce syntactically valid output, or 2) produces an output that is provably false (since its negation is provably true).

4. Benchmark Construction & Tool Implementation

We evaluate our autoformalisation pipeline using a dataset derived from ProcessBench (Zheng et al., 2025). Our approach transforms linear Chain-of-Thought (CoT) reasoning into structured dependency graphs, allowing us to isolate and verify individual reasoning steps while preserving their necessary logical context. We present our brief pipeline in Figure 1.

4.1. Dataset Construction and Pre-processing

ProcessBench (Zheng et al., 2025) consists of natural language reasoning chains of the form shown in Figure 1. As the first step in our pipeline, we divide each of these reasoning chains into a number of "mini-CoTs". Intuitively, our goal here is for each mini-CoTs to encode semantically a single formal statement as in (3). We perform this step because many current AF systems (including the examples we consider below) are designed to output a single statement of this form, rather than a full sequence as in (1), and so we would like the data we consider to match this.

Automated Graph Construction To obtain these mini-CoTs, we first use an LLM to convert each raw CoT in ProcessBench into a DAG format. As illustrated in Figure 1, the model parses the input text, identifies discrete reasoning steps S_i, and defines edges based on logical dependencies. We then iterate over the nodes of this graph to extract individual logical statements that we use as the data points for our benchmark. Crucially, this approach retains the relevant dependencies for every node, ensuring each statement

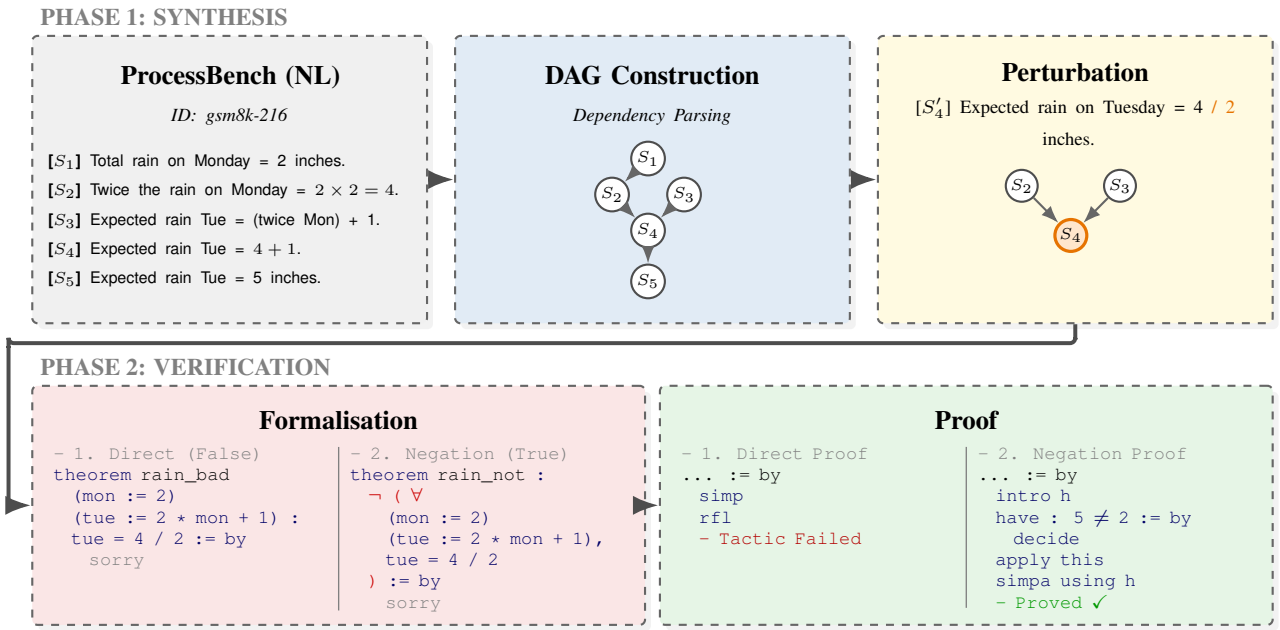

Figure 1. An LLM parses ProcessBenchs natural language steps into a DAG. We then perturb and formalise the nodes associated with the parents into a pair of opposing statements: a direct one and its negation. The prover will try to show both statements and compile the proofs; if the formalisation preserves the invalidity, the compilation will support the negation.

S_i is presented to the autoformalisation pipeline with the context required for its translation and verification. We refer the reader to Appendix B for implementation details and evaluation of this and the next pre-processing stages.

Statement Classification and Filtering Not all nodes in a reasoning chain are suitable for formal verification. During the DAG construction, the LLM concurrently classifies each statement S_i into one of two sets:

1. **Declarative/Unverifiable ($\mathcal{D}$):** Statements that are axiomatic, declarative, or definitions. These are by definition unverifiable in our pipeline.
2. **Verifiable ($\mathcal{O}$):** Statements involving calculations, logical deductions, or derived propositions.

Statements in $\mathcal{D}$ are excluded from the target list of steps to be formalised; however, they remain in the graph to serve as context (predecessors) for dependent statements. Only statements in $\mathcal{O}$ undergo active formalisation and verification.

Our initial source consists of reasoning chains from four datasets of increasing difficulty: GSM8K (Cobbe et al., 2021), MATH (Hendrycks et al., 2021), OlympiadBench (He et al., 2024), and Omni-MATH (Gao et al., 2024a). Following the pre-processing and filtering pipeline described above, we obtained 1,179 processed reasoning chains. From these, we extracted **12,784** individual, verifiable statements (excluding set $\mathcal{D}$).

Perturbation As described in Section 3, a key part of our methodology is a *perturbation step*. For this, we apply a perturbation function Perturb over our dataset as follows:

$$S_i \xmapsto{\text{Perturb}} S_i^{\text{perturbed}}$$

We instantiate Perturb concretely as $\text{Perturb}_{\text{regex}}$ and $\text{Perturb}_{\text{LLM}}$, a Regex-based and an LLM-based approach respectively.

$\text{Perturb}_{\text{regex}}$: this perturbator is the default used in our main experiments, consisting of a series of cascading rules to systematically change input values in the input statement. Since this rule-based approach might occasionally lead to ineffective perturbations, we additionally implement a boolean-valued *Ineffective Perturbation Finder* IPF, and exclude each input S_i where $\text{IPF}(S_i^{\text{perturbed}}) = \text{True}$ from our analysis. Details about implementations of $\text{Perturb}_{\text{regex}}$ and IPF, with ablation study, are included in Appendix C.

$\text{Perturb}_{\text{LLM}}$: In order to evaluate the robustness of our method to different perturbation strategies, we also implement an alternative LLM-based perturbation strategy. Details of prompt used and additional experiments with alternative perturbation strategies are available in appendix G.

Our final benchmark comprises 11697 unperturbed statements, 11697 Regex-perturbed statements (simply referred to as "perturbed" in the next sections) and 11697 LLM-perturbed ones. We refer to this benchmark as Faithform-Bench, with more information available in Appendix E and

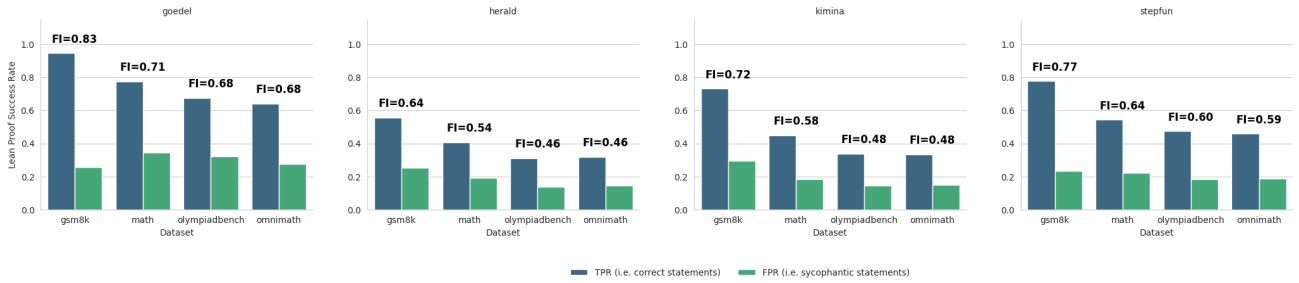

Figure 2. TPR, FPR and FI measures for the SOTA FA methods used across the 4 different datasets using the Regex perturbation strategy.

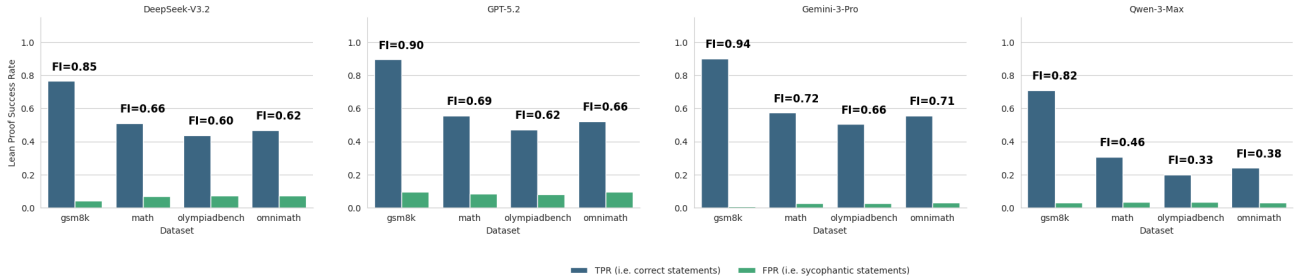

Figure 3. TPR, FPR and FI measures for the End-to-End methods used across the 4 different datasets using the Regex perturbation strategy.

we publicly release it for further exploration.

4.2. Lean Proof Creation

Autoformalisation We employ several compact autoformalisers from the literature—Goedel (Lin et al., 2025) (8B), Herald (Gao et al., 2024b) (7B), Kimina (Wang et al., 2025a) (8B), and StepFun (Wu et al., 2025) (7B). Our pipeline proceeds in two steps. First, we extract the final theorem from the raw output; failures at this stage are flagged as "AF-Fail". Second, we attempt to compile the extracted result as a *by-sorry* theorem. A successful compilation yields a "pass" verdict and advances to the next stage. Failures here are flagged as "TC-Fail".

Proving We attempt to prove both the statement and its negation using DeepSeek-Prover-V2 (Ren et al., 2025). To prevent the prover from sycophancy, where it proves a different statement entirely, we verify that the prover's output matches the formaliser's input. Mismatches are immediately discarded as "Unknown". Valid statements are evaluated at pass@2, resulting in one of three final classifications: "Proved" (direct statement verified), "Refuted" (negation verified), or "Unknown" (neither verified).

5. Experimental Results

5.1. Autoformaliser Results

Figure 2 shows the TPR, FPR, and FI metrics for the perturbed and unperturbed splits of FaithformBench.

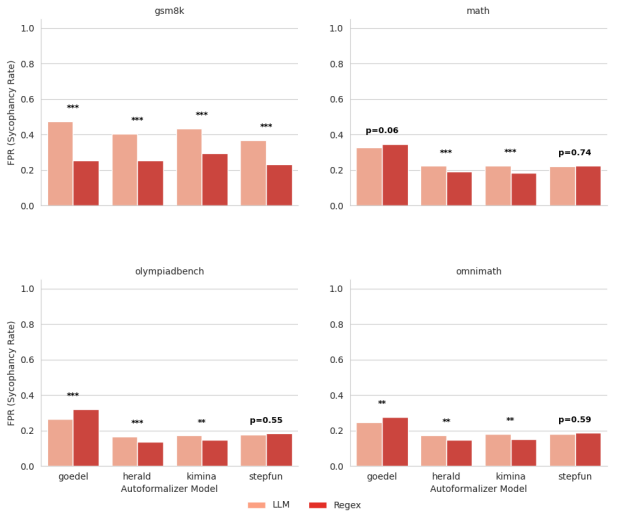

Figure 4. Comparison of different types of perturbations (LLM vs Regex-based) in terms of FPR (i.e. Sycophancy Rate). The asterisks indicate the statistical significance of the difference between the two perturbation methods (Two-Proportion Z-Test). Key: *** ($p < 0.001$), ** ($p < 0.01$), * ($p < 0.05$); non-significant results are labeled with their exact p-value.

TPR AFs show a downward trend of TPR, indicating that they increasingly struggle to create provable Lean statements as the mathematical complexity of the CoTs rises.

FPR All AFs show a remarkable level of sycophancy (this time more preponderant in the easiest GSM8K).

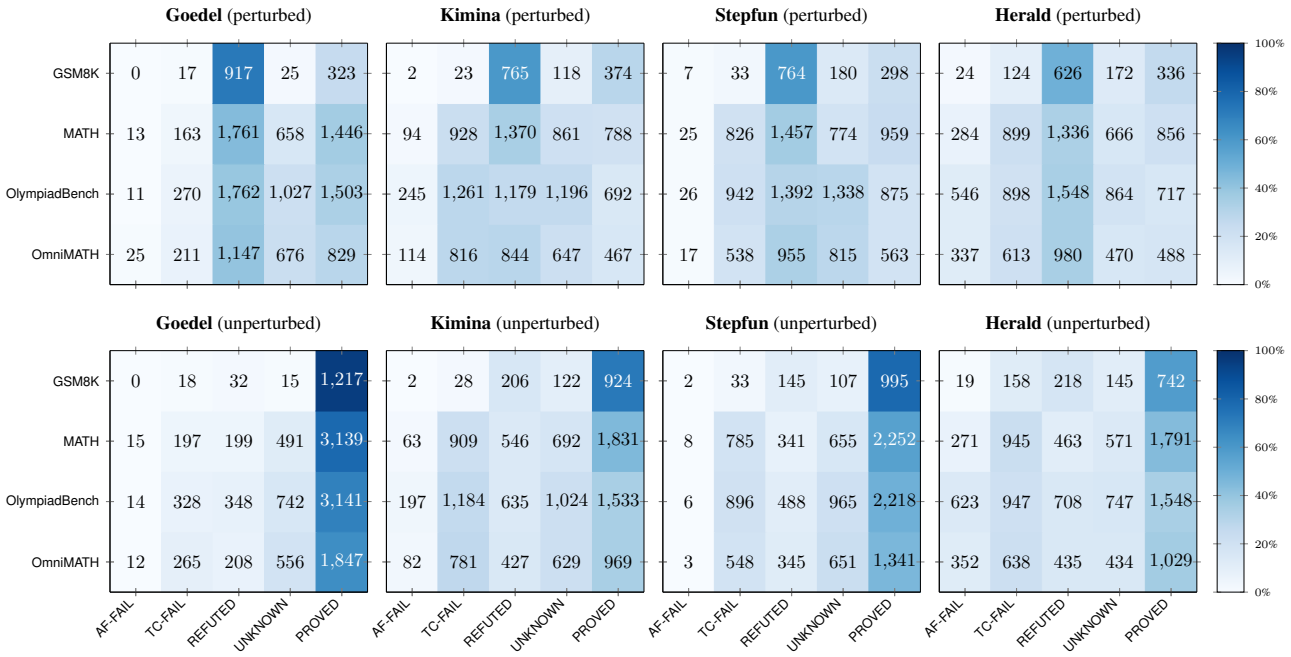

Table 1. Detailed results of 4 AFs on the 4 datasets used.

Goedel, in particular, is the most sycophantic model.

Overall Ranking (FI) Goedel is the most faithful AF in terms of FI for all datasets, followed by Stepfun, Kimina and Herald. Notice also the positive correlation between TPR (validity preservation) and FPR (sycophancy). Goedel owes its high FI in the ranking to high TPRs, which compensate for high FPRs.

Detailed Results Figure 1 breaks down our results in the 5 possible outcomes previously defined. AF-FAIL and TC-FAIL appear low throughout the experiments, while the REFUTED and PROVED cases highlight similar patterns described above. Interestingly, Goedel shows the least REFUTED cases, possibly highlighting problems in autoformalising negations.

5.2. End-to-end Results

Figure 3 illustrates TPR, FPR and FI measures using four frontier models to transform CoTs directly to Lean proofs. The most striking feature of these results is that all models exhibit low levels of sycophancy when compared to AFs (see Figure 2). At the same time, these general LLMs have lower TPRs than the best AF, Goedel, demonstrating that at least the best AFs are indeed better at what they were fine-tuned for, i.e., to translate (correct) natural language statements into *correct* Lean representations, at the cost of exhibiting a higher degree of sycophantic behaviour. In terms of FI, Gemini-3-Pro is the strongest; whereas Qwen-3-Max is the weakest by some distance.

5.3. Comparison of Different Perturbation Strategies

Figure 4 compares the two perturbation strategies: (default) Regex-based vs LLM-based. We see that the perturbation strategy has some effect on the number of sycophantic instances, which is most noticeable in GSM8K and much less relevant in more difficult datasets. The test statistics annotated in the plot further suggest that in certain cases (notably OlympiadBench and OmniMATH) the difference in FPR is not relevant. Overall, LLM-based perturbations lead to more sycophantic cases in the case of GSM8K, while there is no clear trend for the other datasets. The overall ranking is unchanged, with Goedel largely remaining the most sycophantic AF, followed by Stepfun, Kimina and Herald.

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.

References

Achim, T., Best, A., Bietti, A., Der, K., Fédérico, M., Gukov, S., Halpern-Leistner, D., Henningsgard, K., Kudryashov, Y., Meiburg, A., Michelsen, M., Patter-

son, R., Rodriguez, E., Scharff, L., Shanker, V., Sicca, V., Sowrirajan, H., Swope, A., Tamas, M., Tenev, V., Thomm, J., Williams, H., and Wu, L. Aristotle: Imo-level automated theorem proving, 2025. URL https://arxiv.org/abs/2510.01346.

Chen, J., Chen, W., Du, J., Hu, J., Jiang, Z., Jie, A., Jin, X., Jin, X., Li, C., Shi, W., Wang, Z., Wang, M., Wei, C., Wei, S., Xin, H., Yang, F., Gao, W., Yuan, Z., Zhan, T., Zheng, Z., Zhou, T., and Zhu, T. H. Seed-prover 1.5: Mastering undergraduate-level theorem proving via learning from experience, 2025. URL https://arxiv.org/abs/2512.17260.

Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., Hesse, C., and Schulman, J. Training verifiers to solve math word problems. *CoRR*, abs/2110.14168, 2021. URL https://arxiv.org/abs/2110.14168.

Coquand, T. and Huet, G. The Calculus of Constructions. *Information and Computation*, 76(2–3):95–120, 1988. URL https://doi.org/10.1016/0890-5401(88)90005-3.

de Moura, L. and Ullrich, S. The Lean 4 Theorem Prover and Programming Language. In *Automated Deduction – CADE 28*, volume 12699 of *Lecture Notes in Computer Science*, pp. 625–635. Springer, 2021. URL https://doi.org/10.1007/978-3-030-79876-5_37.

Gao, B., Song, F., Yang, Z., Cai, Z., Miao, Y., Dong, Q., Li, L., Ma, C., Chen, L., Xu, R., Tang, Z., Wang, B., Zan, D., Quan, S., Zhang, G., Sha, L., Zhang, Y., Ren, X., Liu, T., and Chang, B. Omni-math: A universal olympiad level mathematic benchmark for large language models. *ArXiv*, abs/2410.07985, 2024a. URL https://api.semanticscholar.org/CorpusID:273233571.

Gao, G., Wang, Y., Jiang, J., Gao, Q., Qin, Z., Xu, T., and Dong, B. Herald: A natural language annotated lean 4 dataset. *CoRR*, abs/2410.10878, 2024b. doi: 10.48550/ARXIV.2410.10878. URL https://doi.org/10.48550/arXiv.2410.10878.

He, C., Luo, R., Bai, Y., Hu, S., Thai, Z. L., Shen, J., Hu, J., Han, X., Huang, Y., Zhang, Y., Liu, J., Qi, L., Liu, Z., and Sun, M. Olympiadbench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. In Ku, L., Martins, A., and Srikumar, V. (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2024, Bangkok, Thailand, August 11-16, 2024*, pp. 3828–3850. Association for Computational Linguistics, 2024. doi: 10.18653/V1/2024.ACL-LONG.211. URL https://doi.org/10.18653/v1/2024.acl-long.211.

Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., and Steinhardt, J. Measuring mathematical problem solving with the MATH dataset. *CoRR*, abs/2103.03874, 2021. URL https://arxiv.org/abs/2103.03874.

Jana, P., Kale, K., Tanriverdi, A. E., Song, C., Vishwanath, S., and Ganesh, V. Proofbridge: Auto-formalization of natural language proofs in lean via joint embeddings. *CoRR*, abs/2510.15681, 2025. doi: 10.48550/ARXIV.2510.15681. URL https://doi.org/10.48550/arXiv.2510.15681.

Jiang, A. Q., Welleck, S., Zhou, J. P., Lacroix, T., Liu, J., Li, W., Jamnik, M., Lample, G., and Wu, Y. Draft, sketch, and prove: Guiding formal theorem provers with informal proofs. In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net, 2023. URL https://openreview.net/forum?id=SMa9EAovKMC.

Kirichenko, P., Ibrahim, M., Chaudhuri, K., and Bell, S. J. Abstentionbench: Reasoning llms fail on unanswerable questions, 2025. URL https://arxiv.org/abs/2506.09038.

Lee, J., Mukherjee, S., Hakkani-Tur, D., and Hockenmaier, J. Reasoningflow: Semantic structure of complex reasoning traces, 2025. URL https://arxiv.org/abs/2506.02532.

Lin, Y., Tang, S., Lyu, B., Yang, Z., Chung, J., Zhao, H., Jiang, L., Geng, Y., Ge, J., Sun, J., Wu, J., Gesi, J., Lu, X., Acuna, D., Yang, K., Lin, H., Choi, Y., Chen, D., Arora, S., and Jin, C. Goedel-prover-v2: Scaling formal theorem proving with scaffolded data synthesis and self-correction. *CoRR*, abs/2508.03613, 2025. doi: 10.48550/ARXIV.2508.03613. URL https://doi.org/10.48550/arXiv.2508.03613.

Liu, C., Yuan, Y., Yin, Y., Xu, Y., Xu, X., Chen, Z., Wang, Y., Shang, L., Liu, Q., and Zhang, M. Safe: Enhancing mathematical reasoning in large language models via retrospective step-aware formal verification. In Che, W., Nabende, J., Shutova, E., and Pilehvar, M. T. (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2025, Vienna, Austria, July 27 - August 1, 2025*, pp. 12171–12186. Association for Computational Linguistics, 2025a. URL https://aclanthology.org/2025.acl-long.594/.

Liu, Q., Zheng, X., Lu, X., Cao, Q., and Yan, J. Rethinking and improving autoformalization: Towards a faithful metric and a dependency retrieval-based approach. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025*. OpenReview.net, 2025b. URL https://openreview.net/forum?id=hUb2At2DsQ.

Liu, Y., Zhu, T., Liu, X., Chen, Y., ZhaoXuan, L., Zhang, J., Bao, K., Luo, T., et al. Generalized tree edit distance (gted): A faithful evaluation metric for statement autoformalization. In *2nd AI for Math Workshop @ ICML 2025*, 2025c.

Liu, Y., Liu, X., Sun, Z., and Hu, W. Answering the unanswerable is to err knowingly: Analyzing and mitigating abstention failures in large reasoning models, 2026. URL https://arxiv.org/abs/2508.18760.

Luong, T. and Lockhart, E. Advanced version of gemini with deep think officially achieves gold-medal standard at the international mathematical olympiad. DeepMind Blog. URL https://deepmind.google/blog/advanced-version-of-gemini-with-deep-think-officially-achieves-gold-medal-standard-at-the-international-mathematical-olympiad/.

Paulson, L. C. Isabelle: The Next 700 Theorem Provers. In Odifreddi, P. (ed.), *Logic and Computer Science*, pp. 361–386. Academic Press, London, 1990.

Perez, E., Ringer, S., Lukosiute, K., Nguyen, K., Chen, E., Heiner, S., Pettit, C., Olsson, C., Kundu, S., Kadavath, S., Jones, A., Chen, A., Mann, B., Israel, B., Seethor, B., McKinnon, C., Olah, C., Yan, D., Amodei, D., Amodei, D., Drain, D., Li, D., Tran-Johnson, E., Khundadze, G., Kernion, J., Landis, J., Kerr, J., Mueller, J., Hyun, J., Landau, J., Ndousse, K., Goldberg, L., Lovitt, L., Lucas, M., Sellitto, M., Zhang, M., Kingsland, N., Elhage, N., Joseph, N., Mercado, N., DasSarma, N., Rausch, O., Larson, R., McCandlish, S., Johnston, S., Kravec, S., El Showk, S., Lanham, T., Telleen-Lawton, T., Brown, T., Henighan, T., Hume, T., Bai, Y., Hatfield-Dodds, Z., Clark, J., Bowman, S. R., Askell, A., Grosse, R., Hernandez, D., Ganguli, D., Hubinger, E., Schiefer, N., and Kaplan, J. Discovering language model behaviors with model-written evaluations. In Rogers, A., Boyd-Graber, J., and Okazaki, N. (eds.), *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 13387–13434, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-acl.847. URL https://aclanthology.org/2023.findings-acl.847/.

Petrov, I., Dekoninck, J., and Vechev, M. T. Brokenmath: A benchmark for sycophancy in theorem proving with llms. *CoRR*, abs/2510.04721, 2025. doi: 10.48550/ARXIV.2510.04721. URL https://doi.org/10.48550/arXiv.2510.04721.

Quan, X., Valentino, M., Dennis, L. A., and Freitas, A. Faithful and robust LLM-driven theorem proving for NLI explanations. In Che, W., Nabende, J., Shutova, E., and Pilehvar, M. T. (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.867.

Ren, Z. Z., Shao, Z., Song, J., Xin, H., Wang, H., Zhao, W., Zhang, L., Fu, Z., Zhu, Q., Yang, D., Wu, Z. F., Gou, Z., Ma, S., Tang, H., Liu, Y., Gao, W., Guo, D., and Ruan, C. Deepseek-prover-v2: Advancing formal mathematical reasoning via reinforcement learning for subgoal decomposition. *CoRR*, abs/2504.21801, 2025. doi: 10.48550/ARXIV.2504.21801. URL https://doi.org/10.48550/arXiv.2504.21801.

Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askell, A., Bowman, S. R., Durmus, E., Hatfield-Dodds, Z., Johnston, S. R., Kravec, S., Maxwell, T., McCandlish, S., Ndousse, K., Rausch, O., Schiefer, N., Yan, D., Zhang, M., and Perez, E. Towards understanding sycophancy in language models. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net, 2024. URL https://openreview.net/forum?id=tvhaxkMKAn.

Tsoukalas, G., Lee, J., Jennings, J., Xin, J., Ding, M., Jennings, M., Thakur, A., and Chaudhuri, S. Putnambench: Evaluating neural theorem-provers on the putnam mathematical competition. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024.

Wang, H., Xin, H., Zheng, C., Liu, Z., Cao, Q., Huang, Y., Xiong, J., Shi, H., Xie, E., Yin, J., Li, Z., and Liang, X. LEGO-prover: Neural theorem proving with growing libraries. In *The Twelfth International Conference on Learning Representations*, 2024a.

Wang, H., Unsal, M., Lin, X., Baksys, M., Liu, J., Santos, M. D., Sung, F., Vinyes, M., Ying, Z., Zhu, Z., Lu, J., Saxcé, H. d., Bailey, B., Song, C., Xiao, C., Zhang, D., Zhang, E., Pu, F., Zhu, H., Liu, J., Bayer, J., Michel, J., Yu, L., Dreyfus-Schmidt, L., Tunstall, L., Pagani, L., Machado, M., Bourigault, P., Wang, R., Polu, S., Barroyer, T., Li, W.-D., Niu, Y., Fleureau, Y., Hu, Y., Yu, Z., Wang, Z., Yang, Z., Liu, Z., and Li, J. Kimina-prover preview: Towards large formal reasoning models with reinforcement learning. 2025a. URL http://arxiv.org/abs/2504.11354.

Wang, H., Unsal, M., Lin, X., Baksys, M., Liu, J., Santos, M. D., Sung, F., Vinyes, M., Ying, Z., Zhu, Z., Lu, J., Saxcé, H. d., Bailey, B., Song, C., Xiao, C., Zhang, D., Zhang, E., Pu, F., Zhu, H., Liu, J., Bayer, J., Michel, J., Yu, L., Dreyfus-Schmidt, L., Tunstall, L., Pagani, L., Machado, M., Bourigault, P., Wang, R., Polu, S., Barroyer, T., Li, W.-D., Niu, Y., Fleureau, Y., Hu, Y., Yu, Z., Wang, Z., Yang, Z., Liu, Z., and Li, J. Kimina-prover preview: Towards large formal reasoning models with reinforcement learning. 2025b. URL http://arxiv.org/abs/2504.11354.

Wang, R., Zhang, J., Jia, Y., Pan, R., Diao, S., Pi, R., and Zhang, T. TheoremLlama: Transforming general-purpose LLMs into lean4 experts. In Al-Onaizan, Y., Bansal, M., and Chen, Y.-N. (eds.), *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 11953–11974, Miami, Florida, USA, November 2024b. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.667. URL https://aclanthology.org/2024.emnlp-main.667/.

Wu, Y., Jiang, A. Q., Li, W., et al. Autoformalization with large language models. In *Advances in Neural Information Processing Systems*, volume 35. Curran Associates, Inc., 2022. URL https://dl.acm.org/doi/10.5555/3600270.3602614.

Wu, Y., Huang, D., Wan, R., Peng, Y., Shang, S., Cao, C., Qi, L., Zhang, R., Du, Z., Yang, J., and Hu, X. Stepfun-formalizer: Unlocking the autoformalization potential of llms through knowledge-reasoning fusion. *CoRR*, abs/2508.04440, 2025. doi: 10.48550/ARXIV.2508.04440. URL https://doi.org/10.48550/arXiv.2508.04440.

Xin, H., Ren, Z., Song, J., Shao, Z., Zhao, W., Wang, H., Liu, B., Zhang, L., Lu, X., Du, Q., Gao, W., Zhang, H., Zhu, Q., Yang, D., Gou, Z., Wu, Z., Luo, F., and Ruan, C. Deepseek-prover-v1.5: Harnessing proof assistant feedback for reinforcement learning and monte-carlo tree search. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=I4YAIwrsXa.

Xiong, Z., Cai, Y., Li, Z., and Wang, Y. Mapping the minds of LLMs: A graph-based analysis of reasoning LLMs. In Christodoulopoulos, C., Chakraborty, T., Rose, C., and Peng, V. (eds.), *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 17751–17763, Suzhou, China, November 2025. Association for Computational Linguistics. ISBN 979-8-89176-332-6. doi: 10.18653/v1/2025.emnlp-main.896. URL https://aclanthology.org/2025.emnlp-main.896/.

Xu, Q., Wang, R., Wang, P., Li, H., and Watt, C. A minimalist proof language for neural theorem proving over isabelle/hol, 2025. URL https://arxiv.org/abs/2507.18885.

Xue, B., Zhu, Q., Wang, R., Wang, S., Wang, H., Hu, M., Mi, F., Wang, Y., Shang, L., Liu, Q., and Wong, K.-F. Reliablemath: Benchmark of reliable mathematical reasoning on large language models, 2025. URL https://arxiv.org/abs/2507.03133.

Yang, K., Swope, A., Gu, A., Chalamala, R., Song, P., Yu, S., Godil, S., Prenger, R., and Anandkumar, A. LeanDojo: Theorem proving with retrieval-augmented language models. In *Neural Information Processing Systems (NeurIPS)*, 2023.

Zheng, C., Zhang, Z., Zhang, B., Lin, R., Lu, K., Yu, B., Liu, D., Zhou, J., and Lin, J. Processbench: Identifying process errors in mathematical reasoning. In Che, W., Nabende, J., Shutova, E., and Pilehvar, M. T. (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2025, Vienna, Austria, July 27 - August 1, 2025*, pp. 1009–1024. Association for Computational Linguistics, 2025. URL https://aclanthology.org/2025.acl-long.50/.

Zheng, K., Han, J. M., and Polu, S. minif2f: a cross-system benchmark for formal olympiad-level mathematics. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=9ZPegFuFTFv.

Zhou, J. P., Staats, C., Li, W., Szegedy, C., Weinberger, K. Q., and Wu, Y. Don't trust: Verify - grounding LLM quantitative reasoning with auto-formalisation. In *ICLR*, pp. 1–20, 2024.

Zhou, K. and Zhang, L. Step-wise formal verification for LLM-based mathematical problem solving. *arXiv preprint arXiv:2505.20869*, 2025. URL https://doi.org/10.48550/arXiv.2505.20869.

A. Correlation Analyses

Here we consider a more in depth correlation analysis of two factors affecting AF performance: namely the effect of reasoning step depth over unperturbed autoformalised statement correctness and the already observed effect of validity preservation on invalidity preservation.

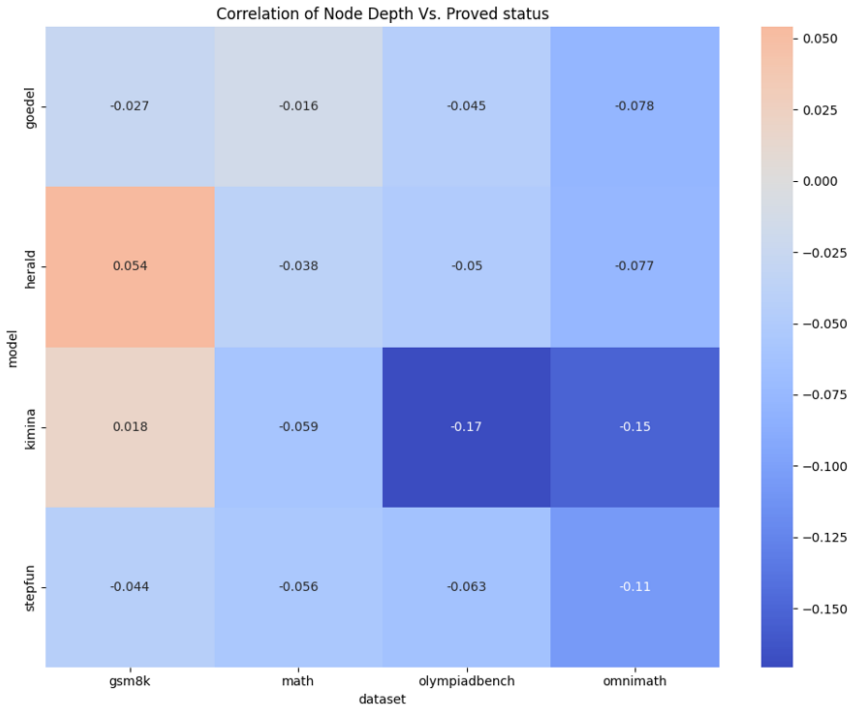

Figure 5. Correlation of number of input statement dependencies against the proved status of the statement in the unperturbed split of FaithformBench.

A.1. Effect of Reasoning Step Depth

In this work we have autoformalised CoT-style reasoning chains produced by LLMs. As mentioned, this choice reflected the growing interest of using AFs for automatic *verification* of AI answers and reasoning from AI models. At the same time, we hypothesize that in fact autoformalising such type of inputs present some unique challenge, such as the difficulty increasing with the depth of the reasoning step (i.e. the index of the input statement relative to which reasoning step is in the originating chain).

By looking at 5 we can indeed confirm that a negative relationship exists between the depth of the unperturbed input reasoning step and the AF being able to correctly prove it. The correlation between the two variables is relatively weak and its sign is not consistent among AFs for GSM8K, the easiest dataset. If we look at other datasets, however, as the relative difficulty of the underlying mathematical problems increase this relationship becomes stronger and consistently negative. Kimina in particular seems to be especially affected by the depth of the input, reflecting a greater struggle in autoformalising statements with longer dependencies. Our work, then, can also help highlight such shortcomings in this new field of application of autoformalisation.

A.2. Effect of AF Validity Preservation on Invalidity Preservation

In the previous section, we have qualitatively observed how AFs like Goedel have a high number of correct Proved statuses for the unpurturbed experiments (i.e. high TPR) but also a high number of incorrect Proved statuses in the perturbed experiments (i.e. high FPR); here we analyse this evidence more in depth to see if indeed a correlation exists between these two variables and what this implies for possible reason behind the observed sycophantic behaviour.

Figure 6 shows the correlation between Proved status of correct input statements under each AF (TPR) and the Proved

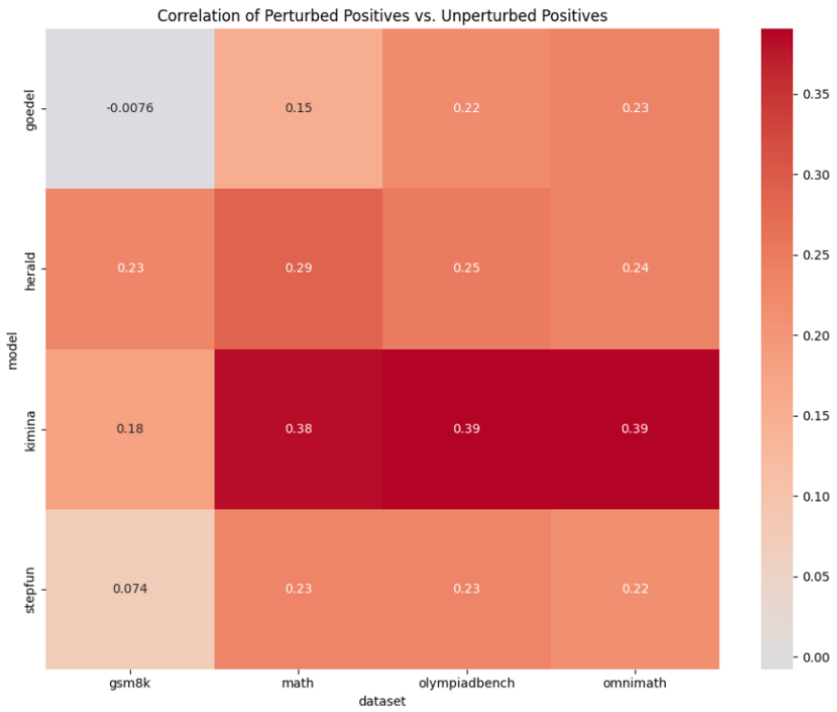

Figure 6. Correlation of proved status of a statement in the perturbed split of FaithformBench against the perturbed status of the same statement in the perturbed split of FaithformBench.

status of the same input statements when perturbed, therefore made incorrect (FPR). From the image is clear that there is a relatively strong correlation between the two variables, meaning that if the AF was able to prove a statement correct when unperturbed than it is also more likely to prove the same statement correct when perturbed, and therefore Positive and Negative Faithfulness appear to be inversely related. From the correlation analysis it is also evident that in fact is Kimina that presents the highest such relation and, therefore, a relatively higher tendency to sycophantic behaviour (as separated by simple AF failures): if, in fact, Goedel appeared to have more absolute sycophantic cases, this can be then partly be explained by its higher success in proving correct statements.

At the same time, this evidence opens up interesting hypotheses about the origin of such sycophantic behaviour. Since we know that an AF that manages to correctly prove the unperturbed statement risk to also silently correct the perturbed counterpart, we can in fact hypothesise that the origin of this behaviour might lay in the classic instruction tuning that these models undertake, which massively emphasises the generation of correct lean proof, instead of possibly incorrect one for incorrect input statements. When the AF "knows" how to correctly formalise the original statement, it would also be more likely to "know" how to correct the perturbed version of it to create a correct lean version of the statement.

More analysis is required to confirm the hypothesis about the origin of the observed sycophantic behaviour. We leave such analysis for future work, among the many other holistic analyses about the behaviour of AFs which the release of our benchmark will enable.

B. Pre-processing

In this part we describe in more depth the various pre-processing steps that we have performed to obtain the final reasoning chains used in our main experiments.

B.1. Filtering of original Datasets

As a first step, we have filtered the various datasets as augmented with reasoning chains and presented in ProcessBench (Zheng et al., 2025), such that:

$$\mathcal{D}_{\text{filtered}} = \{(x, y, r) \in \mathcal{D}_{\text{ProcessBench}} \mid \forall s_i \in r : \text{error}(s_i) = \text{false}\} \tag{5}$$

where x represents the input question, y is the ground-truth answer, $r = (s_1, s_2, \ldots, s_n)$ denotes the reasoning chain, s_i represents the i-th step in the reasoning chain, and $\text{error}(s_i) \in \{\text{true}, \text{false}\}$ indicates whether step s_i contains an error.

B.2. DAG-parsing and Statement Classification

DAG-parsing As mentioned in the main text, we have parsed each reasoning chain into a DAG-like structure in order to obtain the individual steps and the exact dependencies for those steps and reduce noise, as proposed in (Zhou & Zhang, 2025). We follow recent work that parses free-form chain-of-thought traces into structured reasoning graphs that have shown how this task can be easily performed by existing LLMs (Lee et al., 2025; Xiong et al., 2025) and we use a SOTA LLM to perform this step, i.e. GPT 5.2; to ensure the DAG processing was correct, we have prompted GPT 5.2 itself to self-judge the results of the dependencies assignments for each DAG node. Figure 7 confirms that the model is able to perform this passage almost perfectly for every analysed datasets. Even in the cases in which the model might have done an error, it is useful to remember that this might not automatically lead to an error in autoformalisation, since the step might still be valid even if a not fundamental dependency is missing and/or an extra unrelated dependency is included. We therefore conclude that it is safe to assume this preprocessing step will not negatively affect final results significantly.

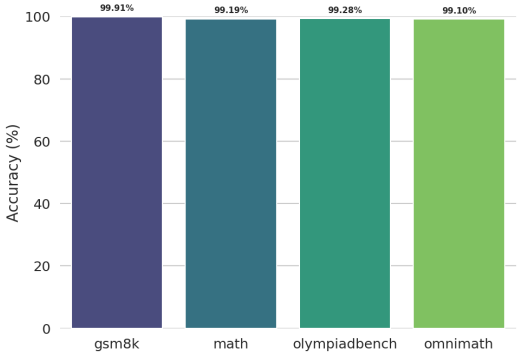

Figure 7. Self-evaluation of DAG parsing step for the 4 datasets using GPT 5.2.

Statement Classification At the same time, we included in the prompt the request to GPT 5.2 to classify each statement s_i into $l_i \in \mathcal{L}_{multi}$, with $\mathcal{L}_{multi}$ including seven statement types defined as following:

1. **Computation:** Statements involving arithmetic or algebraic calculations. These steps perform numerical operations, evaluate expressions, or compute mathematical results. Examples include adding numbers, multiplying terms, evaluating functions, or performing matrix operations.

2. **Logical:** Statements that involve logical inference or deductive reasoning. These include conditional statements (if-then), implications (therefore, thus, hence), logical conclusions drawn from premises, or applications of logical rules and principles.

3. **Declarative:** Statements that assign values to variables, define new variables, or declare mathematical objects. These steps introduce notation, set up the problem space, or establish what symbols represent. Examples include "Let $x = 5$" or "Define $f(x) = x^2 + 1$".

4. **General Knowledge:** Statements that invoke known facts, established theorems, mathematical properties, or domain-specific knowledge. These steps apply existing mathematical results without deriving them, such as citing the Pythagorean theorem, trigonometric identities, or well-known formulas.

5. **Simplification:** Statements that simplify algebraic expressions, reduce fractions, combine like terms, or transform expressions into simpler equivalent forms. These steps make expressions more compact or easier to work with without changing their fundamental value or meaning.

6. **Substitution:** Statements that replace variables or expressions with their equivalent values or alternative representations. These steps involve plugging in known values, replacing variables based on prior definitions, or substituting one expression for another based on established equalities.

7. **Verification:** Statements that check, validate, or confirm the correctness of a result. These steps involve testing solutions, verifying that answers satisfy given conditions, checking work, or confirming that a derived result is consistent with problem constraints.

Once we have classified each statement s_i with the above taxonomy, we derive a simpler rule to create the two sets referenced in 4.1. Based on the nature of declarative assignments being unverifiable and very high frequency in the initial annotation, we partition the statements into a set of "Declarative" statements $\mathcal{D}$ and "Other" statements $\mathcal{O}$ using the indicator function $I(s_i)$:

$$I(s_i) = \begin{cases} \mathcal{D} & \text{if } s_i \text{ is declarative} \\ \mathcal{O} & \text{otherwise} \end{cases}$$

Figure 8 show the results of applying the same self-judging mechanism we have used for verifying the correctness of the DAG dependency parsing, that is we used GPT 5.2 and prompted it to evaluate if each input statement type was correctly classified as belonging to $\mathcal{D}$ or to $\mathcal{O}$. The results clearly show that this step is error free for every dataset analysed.

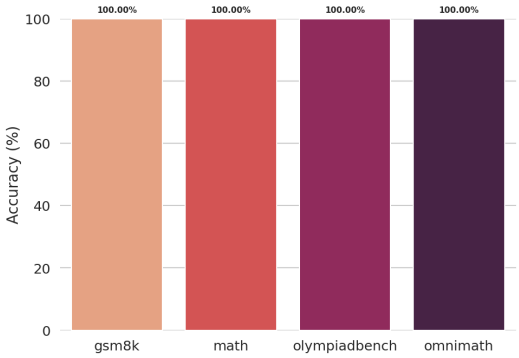

Figure 8. Self-evaluation of statement classification step for the 4 datasets using GPT 5.2.

Finally we also prompt GPT 5.2 to include additional information such as any inference rule that was identified as been represented in the input statement and any note about how the step itself might be verified. Inference rules and notes on verification have not been used in the present work, but they are released together with our datasets (see section E) for future use. Below we include the prompt used in this step:

```
SYSTEM PROMPT:
You are an expert in mathematical logic and proof theory.
Your task is to analyze mathematical reasoning chains and represent them
as directed acyclic graphs (DAGs).

For each step in the reasoning:
1. Classify it as axiom/assumption (no dependencies) or deduction (depends on prior steps)
2. Identify the type of statement (computation, logical inference, etc.)
3. Determine if it's verifiable (can be independently checked)
4. List all dependencies (previous steps it relies on)
5. Name the inference rule used if it's a deduction
```

```
DO NOT INCLUDE THE PROBLEM STATEMENT AMONG THE DAG STEPS.

Be precise and thorough in your analysis.
PROMPT:
Analyze this mathematical reasoning chain and convert it into a structured DAG:

{reasoning_chain}

Create a complete DAG representation with:
- Each reasoning step as a node
- Proper dependency relationships
- Hilbert-style proof annotations
- Verifiability assessment
- Statement type classification
```

C. Experimental Details

C.1. Regex-based Perturbation and Ineffective Perturbation Finder

Regex-based Perturbation We perturb the DAG using a three-level hierarchy of regular-expressionbased transformations. First, we modify the rightmost numeric literal or operator in an equation. As most statements contain equations, this strategy successfully perturbs the majority of the dataset. If no such element is present, we target auxiliary verbs and insert negation (e.g., transforming is larger than into is not larger than). As a final fallback, we prepend the string It is false that to the entire statement. This heuristic may occasionally introduce unintended semantic changesfor example, converting 17 is odd into 13 is odd. We identify and filter out such cases using the method described below.

Ineffective Perturbation Finder To detect ineffective perturbations, we first recover the textual replacement applied to the natural-language statement $S_i^{\text{perturbed}}$ by the regex perturber P_{regex}. We then apply the same replacement to the corresponding unperturbed formalised statement $FL_i^{\text{unperturbed}}$, yielding $FL_i'^{,\text{perturbed}}$ (distinct from $FL_i^{\text{perturbed}}$, which is obtained by applying one of the evaluated AF systems to $S_i^{\text{perturbed}}$). We pass $FL_i'^{,\text{perturbed}}$ to the prover: if it remains provable, we classify the perturbation as ineffective. In this case, we set $\mathsf{IPF}(S_i^{\text{perturbed}}) = \text{True}$ and exclude S_i from FaithformBench.

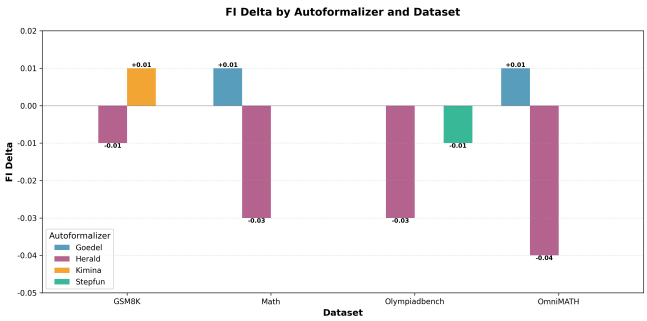

Figure 9. Difference in FI after having excluded examples of Ineffective Perturbations.

As it can be seen from Figure 9, however, the overall effect on the final Faithfulness Index, FI, is minimal and in some cases FI improves after the exclusion of the affected examples, as the percentage of excluded examples was relatively low and often those example comprises instances that led to failure in the unperturbed case.

C.2. Autoformalisation Protocol

Formaliser Parameters We configured all models with a fixed maximum context of 16,384 tokens and set the temperature to zero to minimise stochastic noise. To ensure robustness, we processed the data in batches.

Initial Parsing The pipeline begins by parsing the Lean code generated by the model. We specifically extract the theorem from the final code block and attempt to isolate both its body and its parameters. This step acts as our first filter: if we cannot successfully extract either component, we classify the instance as an AF-Fail and exclude it from downstream processing. Successful instances are then passed to the prover.

C.3. Proving and Verification

Once autoformalisation is complete, we subject the theorems to a two-stage verification process: type-checking and automated proving.

Typechecking First, we compile the extracted theorem using a 'by sorry' tactic to verify its syntactic correctness. Any theorem that fails to compile is flagged as a TC-Fail.

Second, for every valid theorem, we construct a corresponding negation to test for invalidity. We parse the theorem's parameters and body to form a negation statement. If the theorem includes parameters (e.g., $\forall x, P(x)$), we construct the negation as $\neg(\forall x, P(x))$; if it is parameter-free, we simply negate the body.

Finally, we put both the direct statement and its negation to the test. We feed both simultaneously into the DeepSeek Prover V2 (Ren et al., 2025) and evaluate the results at pass@2. This yields three mutually exclusive outcomes:

- **Proved:** The prover verifies the direct statement.
- **Refuted:** The prover verifies the negation statement (indicating the original formalisation was invalid).
- **Unknown:** The prover fails to verify either statement.

C.4. Compute Environment

All experiments were run on a dedicated on-premise server. The system is an ASUS ESC8000A-E13 (non-MGX) equipped with a single AMD EPYC 9555 CPU (64 cores, 3.20,GHz base frequency, up to 4.40,GHz boost, 256,MB L3 cache) and 384,GB of ECC system memory. GPU computation was provided by three NVIDIA H200 NVL Tensor Core GPUs, each with 141,GB of HBM3e memory.

The server runs Ubuntu 22.04 LTS (with limited testing on Ubuntu 24.04 LTS) and uses the NVIDIA CUDA Toolkit and cuDNN. Autoformalizers and provers were executed directly on the host machine using publicly available Hugging Face models.

D. Prompts Details

In our experiments, we have proposed a method based on prompting different SOTA, general LLMs to directly produce the Lean proof without any intermediate step producing a proof-less Lean translation of the input natural language step (as instead we have done for the AFs experiments). In designing the prompt we have used few-shot in-context learning as a prompting technique that has been proven to be effective in formalising informal language to a formal one (Wang et al., 2024b). The full prompt we have used is reported below:

```
# System Instruction:
You are an expert in Lean 4 formalisation. Your task is to
autoformalise informal mathematical problems into Lean 4 code.

Fidelity is paramount: Do NOT correct the user's mathematical
statements. If the input contains a mathematical error
(e.g., "1 + 1 = 3"), you must formalise it exactly as written
(1 + 1 = 3).

Proof Generation: Instead of sorry, append a robust proof tactic
that attempts to solve the problem (e.g., by norm_num,
by linarith, by aesop, by ring, or by dec_trivial).
```

```

Goal: We expect the Lean compiler to throw an error if the
statement is false. Do not try to prevent this error.

# Examples:

## Example Input 1:
CONTEXT: 256 = 2^8. We know that log_2(256) = 8.
STATEMENT: log_2(256) = 8

## Example Output 1:
import Mathlib
import Aesop

set_option maxHeartbeats 0

open BigOperators Real Nat Topology Rat

theorem my_favorite_theorem : Real.logb 2 256 = 8 := by
  norm_num

## Example Input 2:
CONTEXT: Mr. Josue's initial capital was $5000. The first bank
gave him $4000. The second bank gave him $8000.
STATEMENT: 5000 + 4000 + 8000 = 14365

## Example Output 2:
import Mathlib
import Aesop

set_option maxHeartbeats 0

open BigOperators Real Nat Topology Rat

-- Note: This theorem will fail in Lean because the math is incorrect,
-- but the formalisation is faithful to the input.
theorem my_favorite_theorem : 5000 + 4000 + 8000 = 14365 := by
  norm_num

## Example Input 3:
CONTEXT: Solve for x in the expression ((-2)^2 - 11) / (-2 + 1).
STATEMENT: ((-2)^2 - 11) / (-2 + 1) = 7  7 < -1

## Example Output 3:
import Mathlib
import Aesop

set_option maxHeartbeats 0

open BigOperators Real Nat Topology Rat

-- Note: This theorem will fail in Lean because 7 is not less than -1.
theorem my_favorite_theorem :
  ((-2 : )^2 - 11) / (-2 + 1) = 7  (7 : ) < -1 := by

```

```
  norm_num

# Current Task:

CONTEXT: {natural_language_context}
STATEMENT: {natural_language_statement}

Lean 4 Output:
```

E. Benchmarks Details

As part of our contributions, we release the three datasets we produced in performing the experiments described in the main paper in the hope they might foster further research in the area. More details about such datasets are included below:

E.1. Benchmark A: FaithformBench

Description This benchmark evaluates step-level autoformalisation of mathematical chain-of-thought (CoT) reasoning under clean and perturbed conditions.

Task Definition. Each instance corresponds to a single reasoning step extracted from a verified mathematical CoT. Given the natural-language step and associated metadata, models are evaluated on producing a correct formal representation consistent with the proof context.

Data Source. The dataset is derived from PROCESSBENCH, which aggregates CoT solutions for mathematical problems from GSM8K, MATH, OlympiadBench, and OmniMath. ProcessBench augments these problems with reasoning chains generated by QWEN-2.5-INSTRUCT and LLAMA-3-INSTRUCT. While prompting details are not fully specified, the resulting outputs are structured into discrete reasoning steps.

Data Construction. Only reasoning chains verified as correct by ProcessBench are retained. These chains are decomposed into individual steps, each annotated with dependency and inference metadata. From the filtered steps, three subsets are constructed: (i) unperturbed steps, (ii) regex-based perturbations, and (iii) LLM-based perturbations, each comprising approximately one third of the dataset. Perturbations operate at the step level and preserve the original problem identity.

Dataset Statistics. The benchmark contains 45,246 reasoning steps spanning 1,179 unique problems. A summary of the dataset fields is shown in Table 2.

Field	Description
`step_content`	Natural-language reasoning step
`is_perturbed`	Perturbation indicator
`type_of_perturbation`	{unperturbed, regex, LLM}
`problem_id`	Source problem identifier
`node_id`	Step index within CoT
`statement_type`	e.g., computation, deduction
`proof_type`	Predominantly deduction
`dependencies`	References to prior steps
`inference_rule`	Inferrence rule identified in pre-processing (may be missing), e.g. arithmetic
`verification_note`	Notes for verification from pre-processing step

Table 2. Schema for Benchmark A instances.

The majority of steps correspond to computational reasoning, with arithmetic inference rules being the most frequent.

E.2. Benchmark B: Faithform-AF

Goal. This benchmark evaluates the quality and usability of automatically generated Lean statements and proofs corresponding to mathematical chain-of-thought (CoT) reasoning steps. It is intended for training and evaluating LLM-based

Field	Description
`step_content`	Natural-language reasoning step
`is_perturbed`	Perturbation indicator
`type_of_perturbation`	{unperturbed, regex, LLM}
`problem_id`	Source problem identifier
`node_id`	Step index within CoT
`autoformaliser_model`	{Goedel, Kimina, Stepfun, Herald}
`statement_type`	e.g., computation, deduction
`proof_type`	Predominantly deduction
`dependencies`	References to prior steps
`inference_rule`	Inferrence rule identified in pre-processing (may be missing), e.g. arithmetic
`verification_note`	Notes for verification from pre-processing step

Table 3. Schema for Benchmark B instances.

autoformalisers for automatic verification of CoTs.

Task Definition. Each instance augments a reasoning step from Benchmark A with an autoformalised Lean statement and proof. Given the natural-language step and metadata, models are evaluated on producing or verifying Lean formalisations consistent with the mathematical intent of the step.

Data Source. The benchmark is constructed by applying four state-of-the-art autoformalisation models (GOEDEL V2, KIMINA, STEPFUN, and HERALD) to all instances of Benchmark A. For each reasoning step, the corresponding Lean statement and proof produced by each model are recorded.

Data Construction. For every step in Benchmark A, each autoformalisation model generates a Lean formalisation. The resulting statements and proofs are evaluated using the verification protocol described in the main paper, and the outcomes are stored alongside the original step metadata. This produces one instance per *(reasoning step, autoformaliser)* pair, preserving perturbation labels and problem identities from Benchmark A.

Dataset Statistics. The benchmark contains 180,984 instances, corresponding to 4 autoformalisations for each reasoning step. A summary of the dataset schema is provided in Table 3.

Summary Statistics. The dataset spans 1,179 unique problems and 53,378 unique reasoning steps. Each step is paired with four autoformalisations. The majority of instances correspond to computational reasoning with arithmetic inference rules.

Intended Use. This benchmark is designed for fine-tuning and evaluating LLM-based autoformalisers, particularly for automatic verification and robustness analysis of mathematical CoTs.

E.3. Benchmark C: Faithform-End2End

Goal. This benchmark evaluates end-to-end mathematical reasoning by frontier large language models prompted to directly generate Lean proofs from natural-language chain-of-thought (CoT) reasoning steps, without an intermediate autoformalisation or classical prover pipeline.

Task Definition. Each instance corresponds to a reasoning step from Benchmark A paired with a Lean proof produced directly by an end-to-end model. Given the natural-language step and associated metadata, models are evaluated on producing a Lean proof that verifies under the same protocol used for autoformalised proofs.

Data Source. The benchmark is constructed by applying four frontier LLMsGPT-5.2, GEMINI-3-PRO, QWEN-3-MAX, and DEEPSEEK-V3.2to the reasoning steps in Benchmark A. Models are prompted to generate Lean proofs directly; the full prompt specification is provided in Appendix D.

Field	Description
`step_content`	Natural-language reasoning step
`is_perturbed`	Perturbation indicator
`problem_id`	Source problem identifier
`node_id`	Step index within CoT
`end2end_model`	{GPT-5.2, Gemini-3-Pro, Qwen-3-Max, DeepSeek-V3.2}
`statement_type`	e.g., computation, deduction
`proof_type`	Predominantly deduction
`dependencies`	References to prior steps
`inference_rule`	Inferrence rule identified in pre-processing (may be missing)
`verification_note`	Notes for verification from pre-processing step

Table 4. Schema for Benchmark C instances.

Data Construction. For each reasoning step, each end-to-end model produces a Lean proof conditioned on the step content and metadata. Unlike Benchmark B, no explicit autoformalisation or prover-guided routine is used. The original step annotations from the pre-processing procedure described in the main body are included. This results in one instance per *(reasoning step, end-to-end model)* pair, preserving perturbation labels and problem identities.

Dataset Statistics. The benchmark contains 120,656 instances spanning 1,179 unique problems and 27,764 unique reasoning steps. A summary of the dataset schema is shown in Table 4.

Summary Statistics. Each reasoning step is paired with four end-to-end model outputs. As in Benchmarks A and B, the majority of instances correspond to computational or deductive reasoning. Inference rule annotations are occasionally missing due to the absence of explicit autoformalisation.

Intended Use. This benchmark is intended for evaluating and comparing end-to-end proof generation capabilities of frontier LLMs, and for studying trade-offs between modular autoformalisation pipelines and direct proof synthesis.

E.4. Limitations.

Lean formalisations inherit errors and biases from the underlying autoformalisation models. Verification outcomes reflect both reasoning quality and idiosyncrasies of Lean proof search. The benchmark does not enforce global proof coherence across full reasoning chains.

E.5. Reproducibility and Availability.

All dataset construction scripts, perturbation code, and evaluation pipelines will be released publicly after the conclusion of the blind review process.

F. Qualitative Analysis

F.1. Sycophancy Examples

In this section, we present two instances where the autoformalisers exhibit sycophancy. This is not a failure of logic, but a failure of alignment.

EXAMPLE 1: AUTOCORRECT EXAMPLE (STEPFUN)

Problem ID: `gsm8k-390`

In this example, we modify the prompt to claim that $2 \times 54 = 56$ (a clear arithmetic error). The autoformalizer ignores this instruction entirely and proves the corrected result (108).

```
theorem my_favorite_theorem
  (T R O : )
  (h : T = 232)
```

```
(h : R = 54)
(h : O = 2 * R) :
   O = 108 := by
 have h_main : O = 108 := by
 ...
```

EXAMPLE 2: THE ARITHMETIC HALLUCINATION (KIMINA)

Problem ID: `gsm8k-245`

We assert that $1000 - 2 = 686$. Rather than rejecting this obvious falsehood, the formaliser identifies that the input parameters (rates and totals) are mutually exclusive. It proves `False` from the premises and uses that contradiction to close the goal, formally "proving" the arithmetic error.

```
theorem my_favorite_theorem
  (rate_A rate_B : )
  (h_rate_A : rate_A = 5)
  (h_rate_B : rate_B = 2 * rate_A)
  (h_total : 1000 * rate_A + 1000 * rate_B = 686 * rate_A + 686 * rate_B) :
  1000 - 2 = 686 := by
  have h_false : False := by
    have h1 : rate_B = 10 := by
      rw [h_rate_B]
      rw [h_rate_A] ...
```

EXAMPLE 3: *Ex Falso Quodlibet* VIA IMPOSSIBLE CONSTRAINT (KIMINA)

Problem ID: `omnimath-629`

In this example, the perturbed content asserts that $3^{(x+2)} = 51$. The formalizer postulates the formal statement by demonstrating that if x is assumed to be a natural number, no value exists that satisfies the premise (since $3^3 = 27$ and $3^4 = 81$). Consequently, the conclusion follows from a contradiction.

```
theorem my_favorite_theorem (x : ) (h : 3^x = 5) : 3^(x + 2) = 51 := by
   have h : False := by...
```

F.2. Refuted Examples

In these examples, the formaliser fails to prove a statement that is intended to be true in the original context. This failure often stems from the model's inability to bridge the gap between informal mathematical speech and formal rigour. The formaliser translates our example literally, ignoring implicit assumptions such as domain constraints, resulting in a statement that is technically false and thus easily refuted by a counterexample.

EXAMPLE 1: THE MISSING DOMAIN CONSTRAINT (HERALD)

Problem ID: `omnimath-778`

We intend to prove an algebraic identity given $abc = 1$. However, the formaliser translates the variables a, b, c as complex numbers ($\mathbb{C}$) rather than restricting them to positive reals. In the complex domain, the identity does not hold universally. The prover successfully finds a counterexample $(1, -1, 1)$ that refutes the formalised statement.

```
theorem not_a_cub_plus_b_cub_plus_c_cub_eq_3_full :
 ň (
  (a b c : )
 (habc : a * b * c = 1),
 a ^ 3 + b ^ 3 + c ^ 3 = 3
 ) := by
  have h_main : ( (a b c : ) (habc : a * b * c = 1), a ^ 3 + b ^ 3 + c ^ 3 = 3)  False := by
```

```
    intro h
    have h := h (1 : ) (-1 : ) (-1 : )
    have h : (1 : ) * (-1 : ) * (-1 : ) = 1 := by ...
```

EXAMPLE 2: THE FORMULAIC FAILURE (KIMINA)

Problem ID: `olympiadbench-676`

The prompt proposes a specific formula for the GCD of n^3 and $n!$. While the formula may hold for prime numbers, it fails for composite numbers where the power of prime factors grows differently. The formaliser does not attempt to "fix" the conjecture; instead, it tests $n = 4$, identifies the discrepancy, and refutes the theorem.

```
theorem not_my_favorite_theorem_full :
ň (
   {n : } (hn : n  0),
  Nat.gcd (n ^ 3) (n !) =  p in (Nat.primeFactors n),
  p ^ (min (n.factorization p) (n.factorization (p ^ 3)))
) := by
  have h_main : ň ( {n : } (hn : n  0), Nat.gcd (n ^ 3) (n !) =  p ...
```

G. Perturbation Prompt and Sensitivity to Perturbation Prompt

In the main body of this work, we showed that perturbing input reasoning steps using an LLM can lead to different downstream outcomes in certain instances. However, the overall ranking of autoformalisation systems remains stable, and the absolute changes in false positive rate (FPR) and, in particular, failure impact (FI) are generally small and often not statistically significant.

In this appendix, we document the prompts used to generate LLM-based perturbations and present additional analyses examining the sensitivity of downstream FPR to the choice of perturbation prompt. Unless otherwise stated, the *medium* perturbation prompt is used throughout the main paper as the primary LLM-based perturbation condition and as the point of comparison with the regex-based perturbation approach.

G.1. LLM Perturbation Prompts

LLM-based perturbations are generated by instructing the model to modify a single reasoning step while preserving surface fluency and step structure. We define three perturbation levels corresponding to increasing severity of the induced error. Each perturbation level is implemented via a fixed natural-language instruction provided to the LLM.

The exact instructions used are reported below. In all cases, the prompts were fed to DeepSeek V3.2 to produce each perturbed step.

```
level_instructions = {
    "low": (
        "Make a subtle calculation or notational error. "
        "Flip a sign (e.g., + to -), change a single digit in a number, "
        "or make a minor arithmetic slip (e.g., 23 * 2 = 48). "
        "Retain the correct structure/formulas, but corrupt the specific value."
    ),
    "medium": (
        "Make a logical or procedural error. "
        "Apply the wrong formula (e.g., using area formula for circumference), "
        "perform an operation in the wrong order, or draw an invalid inference "
        "from the context. The step should logically fail to follow from the previous steps."
    ),
    "high": (
        "Make a drastic conceptual error or hallucination. "
```

```
        "Introduce mathematical concepts completely unrelated to the problem "
        "(e.g., suddenly using calculus in a simple algebra problem), "
        "invent numbers/variables that were never defined, or state a conclusion "
        "that contradicts basic mathematical axioms."
    )
}
```

Perturbation Semantics. The three perturbation levels are designed to target different failure modes: (i) *low* perturbations introduce localized numerical or symbolic errors while preserving logical structure; (ii) *medium* perturbations break logical validity at the step level while remaining contextually plausible; and (iii) *high* perturbations introduce conceptually invalid or hallucinatory content that is clearly inconsistent with the problem context.

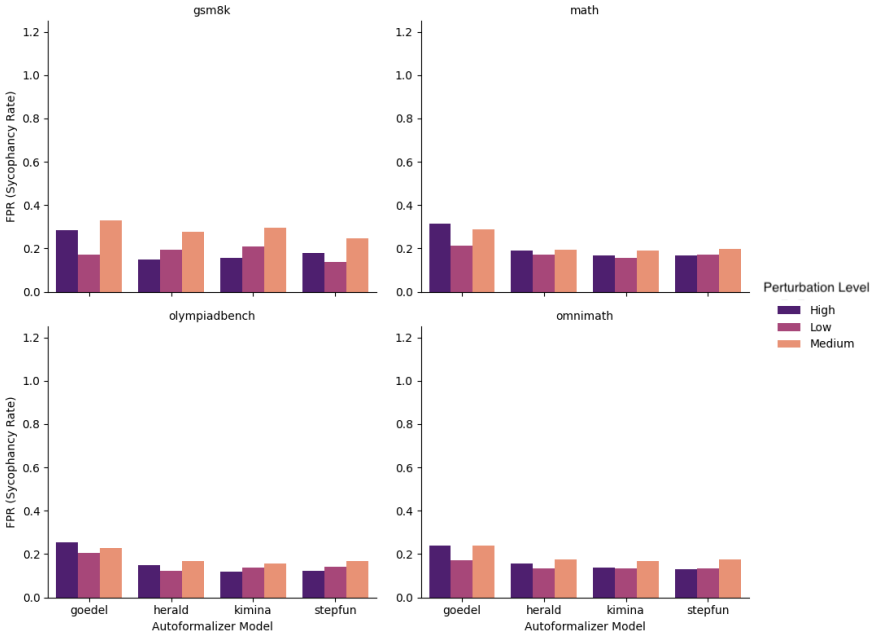

Figure 10. Comparison of FPR (i.e. Sycophantic Rate) obtained with the three different perturbation prompts described above and DeepSeek V3.2 across different datasets and AFs.

G.2. Sensitivity to Perturbation Prompt

Figure 10 reports the false positive rate (FPR) for the four autoformalisation models under three LLM-based perturbation prompts (*low*, *medium*, *high*), stratified by problem source (GSM8K, MATH, OlympiadBench, OmniMath). For each source dataset we have used 50% of the original set used in the main paper.

Across all datasets, we observe a consistent qualitative pattern. First, the relative ranking of autoformalisation models is stable across perturbation levels: GOEDEL exhibits the highest FPR, followed by HERALD, with KIMINA and STEPFUN consistently achieving lower FPR. This ordering holds independently of perturbation severity.

Second, the absolute effect of perturbation severity on FPR is moderate. Moving from *low* to *medium* perturbations generally increases FPR across all models and datasets, while the difference between *medium* and *high* perturbations is smaller and, in several cases, reversed. In particular, for GSM8K and OlympiadBench, *medium* perturbations often induce higher FPR than *high* perturbations, suggesting that logically plausible but incorrect steps are more challenging for autoformalisers than overtly nonsensical ones.

Third, the sensitivity to perturbation prompt varies by dataset. GSM8K exhibits the largest spread in FPR across perturbation levels, while MATH and OmniMath show more compressed ranges. This indicates that perturbation severity interacts with dataset-specific properties, such as problem length and reasoning complexity.

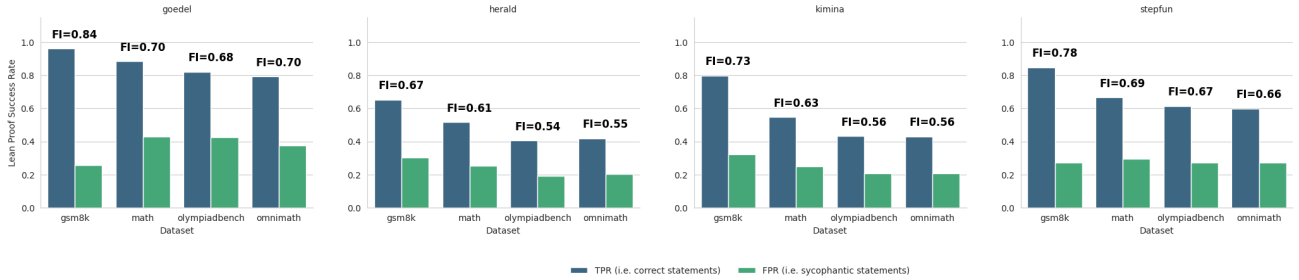

Figure 11. Results obtained with flagging rule $\mathcal{F}_{strict}$ and Regex perturbation for the 4 AFs across the 4 datasets.

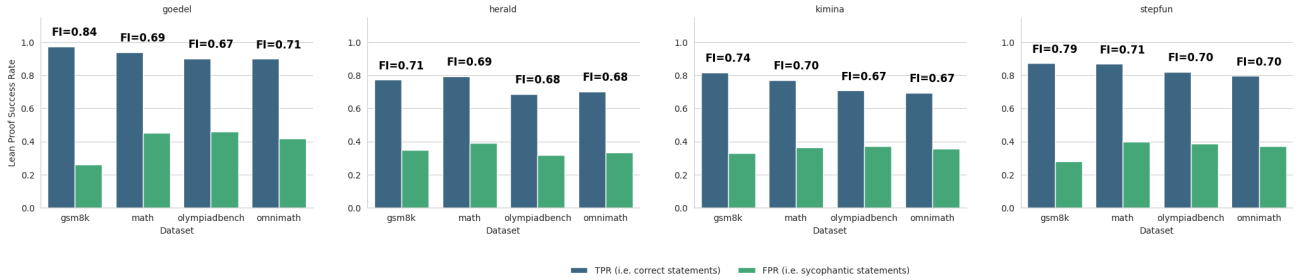

Figure 12. Results obtained with flagging rule $\mathcal{F}_{stricter}$ and Regex perturbation for the 4 AFs across the 4 datasets.

Overall, these results indicate that while perturbation prompt choice affects absolute FPR values, it does not alter the comparative conclusions of the main paper. The *medium* perturbation prompt induces consistently higher or comparable FPR than the *low* prompt and remains close to the worst-case behavior observed under *high* perturbations, making it a reasonable default for stress-testing autoformalisation robustness.

H. Sensitivity to Metric Leniency

Here we consider how results change as we restrict the analysis to the two category of greater interest, namely "Proved" and "Refuted" and as we exclude samples resulting in the other output categories of "AF-Fail", "TC-Fail" and "Unknown" (see section **??** for the definition of the various outcome categories).

In doing so, we first filter out all the results from our experiments that belong to the "Unknown" category and we then recompute the metrics on this new set. Figure 11 shows the resulting metrics under such a change.

We then also consider the case in which we retain only the samples having a "Proved" and "Refuted" output category, therefore excluding also "AF-Fail" and "TC-Fail" outcomes. Figure 12 shows the metrics obtained with this stricter set.

By looking at figure 11, it is evident that the AFs ranking do not appear to change if we just exclude the "Unknown" category. Apart from the case of the easier GSM8K, however, we see that Stepfun become almost as good as Goedel in this scenario, suggesting that most False Positives from this AF were included in the "Unknown" outcome category (which includes cases such as timeouts and sycophantic prover rather than AF).

The change is even more striking if we consider the results in figure 12, where the AF ranking appears dramatically changed with Stepfun now being the best AF and Goedel being comparable to Herald and Kimina, at least for the datasets other than GSM8K. In all these cases, however, the evidence indicate a higher rate of failures on behalf of Kimina and Herald rather than a fundamentally different interpretations of the results under different flagging rules, which can also be confirmed by looking at the higher number of "AF-Fail" and "TC-Fail" under the 2 AFs (see table 1). As the denominator gets smaller in excluding such occurrences, then, the results for these 2 AFs get inflated.

Since we are interested in a holistic evaluation of the AFs, then, in the main text we have considered all the cases including also such failures, as they are still dependent on the model used in auto-formalisation, as highlighted by the differences between different AFs.